

MARCUS RODRIGUEZ

☎ 555-123-4567 ✉ marcus.rodriguez@email.com 🔗 linkedin.com/in/marcus-rodriguez 🐙 github.com/marcus-dev

Education

Massachusetts Institute of Technology <i>Bachelor of Science in Computer Science and Engineering</i>	Sep. 2018 – May 2022 <i>Cambridge, MA</i>
Google AI Research Internship Program <i>Research Focus: Large Language Models and Knowledge Graphs</i>	Jun. 2021 – Aug. 2021 <i>Mountain View, CA</i>

Experience

Senior Software Engineer <i>Anthropic</i> <ul style="list-style-type: none">Developed core infrastructure for Claude's training pipeline, processing 50+ billion tokens dailyImplemented distributed training optimizations reducing model training time by 35%Led team of 6 engineers building safety evaluation frameworks for large language modelsDesigned and deployed real-time monitoring systems handling 1M+ API requests per minute	Jan. 2023 – Present <i>San Francisco, CA</i>
Machine Learning Engineer <i>OpenAI</i> <ul style="list-style-type: none">Built reinforcement learning from human feedback (RLHF) systems for GPT model alignmentOptimized inference serving infrastructure reducing latency by 40% across all model endpointsCollaborated on research leading to 3 publications in top-tier ML conferencesMentored 5 junior engineers and established best practices for model evaluation	Jun. 2022 – Dec. 2022 <i>San Francisco, CA</i>
Software Engineering Intern <i>Meta AI Research</i> <ul style="list-style-type: none">Developed novel attention mechanisms for multimodal transformers processing text and imagesImplemented efficient batching algorithms improving training throughput by 25%Created comprehensive testing suite with 95% code coverage for ML training pipelines	May 2021 – Aug. 2021 <i>Menlo Park, CA</i>

Projects

Neural Code Search Engine <i>Python, PyTorch, Elasticsearch, React</i> <ul style="list-style-type: none">Built semantic code search system using CodeBERT embeddings, indexing 10M+ code repositoriesAchieved 92% accuracy on code similarity tasks, deployed to 50K+ developers at Anthropic	2023
Real-time Sentiment Analysis API <i>FastAPI, Docker, AWS, Redis</i> <ul style="list-style-type: none">Designed scalable microservices architecture processing 100K+ social media posts per hourImplemented custom transformer models fine-tuned on domain-specific data with 96% F1-score	2022
Distributed Training Framework <i>PyTorch, Kubernetes, gRPC</i> <ul style="list-style-type: none">Open-source framework for efficient multi-GPU training, adopted by 15+ research institutionsReduced communication overhead by 60% using gradient compression and asynchronous updates	2021

Technical Skills

Languages: Python, C++, JavaScript, Go, Rust, SQL
Frameworks: PyTorch, TensorFlow, React, FastAPI, Django, Flask
Developer Tools: Git, Docker, Kubernetes, AWS, GCP, Jenkins, Terraform
Libraries: NumPy, Pandas, Scikit-learn, Hugging Face, OpenCV, CUDA
Databases: PostgreSQL, MongoDB, Redis, Elasticsearch, Pinecone

Leadership & Activities

MIT AI Safety Club <i>President and Founder</i> <ul style="list-style-type: none">Founded student organization focused on AI alignment research, growing membership to 200+ studentsOrganized 12 technical workshops and hosted speakers from leading AI research labs	Sep. 2020 – May 2022 <i>Cambridge, MA</i>
Google Summer of Code <i>Open Source Contributor - TensorFlow</i> <ul style="list-style-type: none">Contributed performance optimizations to TensorFlow's distributed training moduleMentored by Google engineers, code merged into main repository used by millions of developers	May 2020 – Aug. 2020 <i>Remote</i>